

Part 1: Research & Documentation

Types of Views in SQL Server

1. Standard View (Regular View)

What is it?

A **Standard View** is a virtual table created using a SELECT statement.

It does **not store data physically**—it only stores the query definition and retrieves data from base tables when queried.

Key Differences

- Data is **not stored** in the view
- Always reflects **current data** from underlying tables
- Simplest and most commonly used view type
- No performance improvement by itself
- **Real-Life Use Case**

University System

- A view that shows student name, course name, and grade by joining multiple tables.
- Used to simplify complex queries for lecturers or admins.

Example:

```
CREATE VIEW vw_StudentResults AS
SELECT s.StudentName, c.CourseName, r.Grade
FROM Students s
JOIN Results r ON s.StudentID = r.StudentID
JOIN Courses c ON r.CourseID = c.CourseID;
```

Limitations & Performance Considerations

- No data storage → query runs every time
- Performance depends on underlying tables and joins

- Cannot use indexes directly on the view
- Heavy queries may be slow on large datasets

2. Indexed View (Materialized View)

What is it?

An **Indexed View** stores the result set **physically on disk** by creating a **clustered index** on the view.

SQL Server maintains the view data automatically.

Key Differences

- Data is stored physically
- Improves performance for complex aggregations
- Requires SCHEMABINDING
- Automatically updated when base tables change

Real-Life Use Case

Banking System

- A view that calculates **total balance per branch**
- Used frequently in reports and dashboards

Example:

```
CREATE VIEW vw_BranchBalance
WITH SCHEMABINDING
AS
SELECT BranchID, SUM(Balance) AS TotalBalance
FROM dbo.Accounts
GROUP BY BranchID;
```

```
CREATE UNIQUE CLUSTERED INDEX idx_BranchBalance
ON vw_BranchBalance (BranchID);
```

Limitations & Performance Considerations

- Slower INSERT, UPDATE, DELETE operations
- Strict rules (no SELECT *, no outer joins, no non-deterministic functions)
- Requires additional storage
- Best for **read-heavy systems**, not write-heavy ones

3. Partitioned View (Union View)

What is it?

A **Partitioned View** combines data from **multiple tables** (usually partitioned by range) using UNION ALL.

Each table stores a portion of the data.

Key Differences

- Data stored in **multiple tables**
- Uses UNION ALL
- Useful for **horizontal partitioning**
- Improves performance by scanning only relevant tables

Real-Life Use Case

E-Commerce System

- Orders split by year into separate tables:
 - Orders_2023
 - Orders_2024
- The view allows querying all orders as one table

Example:

```
CREATE VIEW vw_AllOrders AS
SELECT * FROM Orders_2023
UNION ALL
SELECT * FROM Orders_2024;
```

Limitations & Performance Considerations

- Tables must have identical structure
- No automatic partition management
- Complex maintenance as data grows
- Poor design can hurt performance

Which Types of Views Allow DML Operations?

Standard (Simple) Views

- ✓ **Allow DML** if certain conditions are met
- ✓ Most commonly used for INSERT, UPDATE, DELETE

Updatable when:

- View is based on **one base table**
- No aggregation or grouping
- No joins
- No DISTINCT
- No computed columns
- No GROUP BY / HAVING

 SQL Server directly updates the underlying table.

Indexed Views

DML not allowed directly on the view

- You **cannot INSERT/UPDATE/DELETE** data through an indexed view
- You must modify the **base tables**
- SQL Server automatically maintains the indexed view

Partitioned Views (UNION ALL Views)

 Limited DML support

✓ INSERT allowed **only if:**

- The inserted row belongs to **one specific partition**
- A CHECK constraint exists to identify the correct table

 UPDATE/DELETE often restricted and complex

Restrictions and Limitations for DML on Views

DML operations **are NOT allowed** on a view if it contains:

- 🚫 JOINs (multiple tables)
- 🚫 GROUP BY or HAVING
- 🚫 Aggregate functions (SUM, COUNT, AVG...)
- 🚫 DISTINCT
- 🚫 UNION / UNION ALL (except limited INSERT)
- 🚫 Computed or derived columns
- 🚫 Subqueries in SELECT
- 🚫 TOP with ORDER BY

 **Key Rule:**

If SQL Server cannot clearly map a change to **one row in one base table**, DML is blocked.

3 Real-Life Examples of Updating Views

HR System Example

Table: Employees

| EmployeeID | Name | Salary | Department | IsActive |

View: ActiveEmployeesView

```
CREATE VIEW ActiveEmployeesView AS
SELECT EmployeeID, Name, Salary
FROM Employees
WHERE IsActive = 1;
```

✓ HR updates salaries **through the view**:

```
UPDATE ActiveEmployeesView
SET Salary = Salary + 200
WHERE EmployeeID = 101;
```

✗ The update affects the **Employees** table directly.

✓ Benefits:

- Security (HR can't see inactive employees)
- Simpler interface
- Controlled access

E-Commerce Orders Example

View: PendingOrdersView

```
CREATE VIEW PendingOrdersView AS
SELECT OrderID, CustomerID, Status
```

```
FROM Orders  
WHERE Status = 'Pending';
```

✓ Mark an order as shipped:

```
UPDATE PendingOrdersView  
SET Status = 'Shipped'  
WHERE OrderID = 5001;
```

2. How Can Views Simplify Complex Queries?

How Views Simplify JOIN-Heavy Queries

A **view** is a stored SELECT statement that behaves like a virtual table.

Key benefits for complex JOINs:

- **Removes repeated JOIN logic**
Write the JOIN once → reuse it everywhere.
- **Improves readability**
Call center agents or analysts query a simple view instead of a long SQL statement.
- **Reduces errors**
Less copy-paste = fewer JOIN mistakes.
- **Security control**
Users can access the view without touching base tables.

- **Centralized changes**
If table structure changes, update the view—not every query.

2

Banking Scenario: Call Center Account Summary

Business need

Call center agents frequently need to see:

- Customer name
- Account number
- Account type
- Balance
- Account status

Instead of writing JOINS every time, we create a **view**.

3

Example Tables (Simplified)

```
Customers (
    CustomerID INT,
    FullName VARCHAR(100),
    Phone VARCHAR(20)
);
```

```
Accounts (
    AccountID INT,
    CustomerID INT,
    AccountType VARCHAR(20),
    Balance DECIMAL(10,2),
    Status VARCHAR(20)
);
```



JOIN-Heavy Query (Without a View)

```
SELECT
    c.CustomerID,
    c.FullName,
    a.AccountID,
    a.AccountType,
    a.Balance,
    a.Status
FROM Customers c
JOIN Accounts a
    ON c.CustomerID = a.CustomerID;
```

✗ Problem:

This query must be rewritten **every time** an agent or report needs account details.



Creating a View to Simplify the Query

```
CREATE VIEW vw_CustomerAccountSummary
AS
SELECT
    c.CustomerID,
    c.FullName,
    c.Phone,
    a.AccountID,
    a.AccountType,
    a.Balance,
    a.Status
FROM Customers c
JOIN Accounts a
    ON c.CustomerID = a.CustomerID;
```

✓ JOIN logic is now stored once.

6

Using the View (Much Simpler!)

```
SELECT *
FROM vw_CustomerAccountSummary
WHERE Status = 'Active';
```

Or:

```
SELECT FullName, AccountType, Balance
FROM vw_CustomerAccountSummary
WHERE CustomerID = 101;
```



- No JOINS
- Easy to read
- Faster to write
- Perfect for call center dashboards

7

Real-Life Benefit for Call Center Agents

Without View

- Long JOIN queries
- Higher risk of errors
- Hard to maintain
- Requires SQL expertise

With View

- Simple SELECT
- Consistent results
- Easy to update
- Beginner-friendly

8

Summary

- Views simplify complex JOIN-heavy queries
- They act as pre-built query templates

- Ideal for **frequently accessed data** like banking account summaries
- Improve **readability, security, and maintainability**